

```

#Linear Regression Assumption
# Importing Required Libraries
import pandas as pd
import numpy as np
import seaborn as sns
import matplotlib.pyplot as plt
from sklearn.model_selection import train_test_split
from sklearn.linear_model import LinearRegression
from sklearn.metrics import mean_squared_error,
mean_absolute_error, r2_score
from statsmodels.stats.outliers_influence import
variance_inflation_factor
import statsmodels.api as sm

# Load the Data
df = pd.read_csv('Dataset1.csv') # Replace with your
dataset file name

# ... (rest of your data cleaning and preprocessing)

# Define Target Variable and Features
# Update: Include all columns or select specific columns
generated by get_dummies
X = df[['Age', 'Education', 'Earnings_1974',
'Earnings_1975'] + [col for col in df.columns if
col.startswith(('Race_', 'Hisp_', 'MaritalStatus_'))]]

# Convert 'Education' column to numeric representation
using mapping:
education_mapping = {
    'LessThanHighSchool': 0,
    'HighSchool': 1,
    'JuniorCollege': 2,

```

```

    'Bachelor': 3,
    'Graduate': 4 # Add more levels as needed
}
X['Education'] = X['Education'].map(education_mapping)

# Ensure all columns in X are numeric and handle
missing/infinite values
X = X.astype(float) # This line is added to convert all
columns in X to float type.

# Replace infinite values with NaN
X.replace([np.inf, -np.inf], np.nan, inplace=True)

# Impute or remove NaN values
# Option 1: Impute with the mean
# X.fillna(X.mean(), inplace=True)

# Option 2: Remove rows with NaN values
X.dropna(inplace=True)
df = df[df.index.isin(X.index)] # Update df to match the
rows in X after dropping NaNs

y = df['Earnings_1978']

# ... (rest of your code)
# Train-Test Split
X_train, X_test, y_train, y_test = train_test_split(X, y,
test_size=0.2, random_state=42)

# Check Multicollinearity (VIF)
vif_data = pd.DataFrame()
vif_data["Feature"] = X.columns

```

```
vif_data["VIF"] = [variance_inflation_factor(X.values, i)
for i in range(X.shape[1])]
print(vif_data)

# Fit Linear Regression Model
model = LinearRegression()
model.fit(X_train, y_train)

# Coefficients and Intercept
print("Coefficients:", model.coef_)
print("Intercept:", model.intercept_)

# Predictions
y_pred = model.predict(X_test)

# Evaluate Model
mse = mean_squared_error(y_test, y_pred)
mae = mean_absolute_error(y_test, y_pred)
r2 = r2_score(y_test, y_pred)

print(f"Mean Squared Error: {mse}")
print(f"Mean Absolute Error: {mae}")
print(f"R-squared: {r2}")

# Assumptions Checking

# Residuals
residuals = y_test - y_pred

# 1. Linearity
plt.scatter(y_pred, residuals)
plt.axhline(0, color='red', linestyle='--')
```

```
plt.xlabel('Predicted Values')
plt.ylabel('Residuals')
plt.title('Residuals vs Fitted Values (Linearity Check)')
plt.show()

# 2. Homoscedasticity
sns.scatterplot(x=y_pred, y=np.abs(residuals))
plt.axhline(np.mean(np.abs(residuals)), color='red',
linestyle='--')
plt.title("Residuals vs Fitted Values
(Homoscedasticity)")
plt.show()

# 3. Normality of Residuals
sm.qqplot(residuals, line='s')
plt.title("Q-Q Plot (Normality of Residuals)")
plt.show()

# 4. Autocorrelation (Durbin-Watson Test)
dw = sm.stats.durbin_watson(residuals)
print(f"Durbin-Watson Statistic: {dw}")

# Outliers Detection (Cook's Distance)
influence = sm.OLS(y_train,
sm.add_constant(X_train)).fit().get_influence()
(c, _) = influence.cooks_distance
# Remove the 'use_line_collection' argument
plt.stem(np.arange(len(c)), c, markerfmt=",")
plt.title("Cook's Distance for Outlier Detection")
plt.xlabel("Observation Index")
plt.ylabel("Cook's Distance")
plt.show()
```

```
# Final Prediction for 1978 Earnings
final_predictions = model.predict(X)
df['Predicted_Earnings_1978'] = final_predictions

print(df[['Age', 'Education', 'Earnings_1974',
'Earnings_1975', 'Predicted_Earnings_1978']].head())

# Save Results to a File
df.to_csv('predicted_labour_earnings_1978.csv',
index=False)
```